



Ten training runs per arm, evaluated on both test sets



paired difference +2.63 pp, 95% CI [+0.29, +4.98]
 $p = 0.032$ (paired t), 0.041 (Wilcoxon); classifier ahead in 7 of 10 seeds

paired difference +2.79 pp, 95% CI [+1.09, +4.49]
 $p = 0.005$ (paired t), 0.014 (Wilcoxon); classifier ahead in 9 of 10 seeds